# LLM Leaderboard Exploration

| Model | Features | | Intelligence | Price | Speed | Latency | End-to-End Response Time | Further Analysis |
|---|---|---|---|---|---|---|---|---|
| | Context Window | Creator | Artificial Analysis Intelligence Index | Cost per Task USD | Median Tokens/s | Latency First Chunk (s) | Total Response (s) | |
| Claude Opus 5 (max) | 1M | Anthropic | 63 | $2.34 | 56 | 40.66 | 49.65 | Model ↗ Providers ↗ |
| Claude Opus 5 (xhigh) | 1M | Anthropic | 63 | $1.80 | 52 | 29.84 | 39.41 | Model ↗ Providers ↗ |
| Claude Fable 5 (with fallback) | 1M | Anthropic | 62 | $3.14 | 69 | 97.83 | 105.03 | Model ↗ Providers ↗ |
| Claude Opus 5 (high) | 1M | Anthropic | 61 | $1.23 | 53 | 15.85 | 25.37 | Model ↗ Providers ↗ |
| GPT-5.6 Sol (max) | 1M | OpenAI | 61 | $1.23 | 75 | 104.85 | 111.54 | Model ↗ Providers ↗ |
| Grok 4.6 (high) | 500k | SpaceXAI | 61 | $0.84 | 67 | 43.66 | 51.15 | Model ↗ Providers ↗ |
| Grok 4.6 (xhigh) | 500k | SpaceXAI | 60 | $1.04 | 65 | 46.08 | 53.80 | Model ↗ Providers ↗ |
| Kimi K3 (max) | 1.05M | Kimi | 60 | $0.84 | 35 | 3.30 | 74.48 | Model ↗ Providers ↗ |
| GLM-5.3 (max) | 1M | Z AI | 60 | $0.68 | -- | -- | -- | Model ↗ Providers ↗ |

| Model | Features | | Intelligence | Price | Speed | Latency | End-to-End Response Time | Further Analysis |
|---|---|---|---|---|---|---|---|---|
| | Context Window | Creator | Artificial Analysis Intelligence Index | Cost per Task USD | Median Tokens/s | Latency First Chunk (s) | Total Response (s) | |
| Claude Sonnet 5 (Non-reasoning) | 1M | Anthropic | 43 | $0.42 | 67 | 1.84 | 9.32 | Model ↗ Providers ↗ |

| Model | Features | | Intelligence | Price | Speed | Latency | End-to-End Response Time | Further Analysis |
|---|---|---|---|---|---|---|---|---|
| | Context Window | Creator | Artificial Analysis Intelligence Index | Cost per Task USD | Median Tokens/s | Latency First Chunk (s) | Total Response (s) | |
| GPT-5.6 Luna (Non-reasoning) | 1M | OpenAI | 27 | $0.01 | 137 | 0.77 | 4.42 | Model ↗ Providers ↗ |

# Gemini 3.6 Flash

| Model | Features | | Intelligence | Price | Speed | Latency | End-to-End Response Time | Further Analysis |
|---|---|---|---|---|---|---|---|---|
| | Context Window | Creator | Artificial Analysis Intelligence Index | Cost per Task USD | Median Tokens/s | Latency First Chunk (s) | Total Response (s) | |
| **Gemini 3.6 Flash** | 1M | Google | 52 | $0.34 | 209 | 15.66 | 18.06 | Model ↗ \| Providers ↗ |

Act as an expert Machine Learning lecturer and tutor. Teach me Supervised Learning from absolute beginner level to an intermediate/advanced understanding.

Welcome to **Supervised Machine Learning**. Think of supervised learning as learning with a **teacher**. You feed the algorithm training examples where both the input features (the data) and the correct labels (the answers) are provided. The model's job is to find the mathematical patterns connecting inputs to outputs so it can accurately predict answers for unseen data.

**1. The Foundation: Problem Types**

All supervised learning problems fall into two primary categories based on the target label ($y$):

- **Regression (Continuous Target):** Predicting a continuous numerical value.
  - *Examples:* House price, temperature, stock price.
- **Classification (Discrete Target):** Predicting a categorical label.
  - *Binary:* Default/No Default (0 or 1).
  - *Multiclass:* Dog, Cat, or Bird (0, 1, or 2).

Ask Gemini

Flash

Gemini is AI and can make mistakes.

# Qwen3.7 plus

| Model | Features | | Intelligence | Price | Speed | Latency | End-to-End Response Time | Further Analysis |
|---|---|---|---|---|---|---|---|---|
| | Context Window | Creator | Artificial Analysis Intelligence Index | Cost per Task USD | Median Tokens/s | Latency First Chunk (s) | Total Response (s) | |
| **Qwen3.7 Plus** | 1M | Alibaba | 39 | $0.24 | 56 | 2.08 | 46.59 | Model ↗ Providers ↗ |

Act as an expert Machine Learning lecturer and tutor. Teach me Supervised Learning from absolute beginner level to an intermediate/advanced understanding.

Thinking completed

Welcome to class! I'm thrilled to be your tutor on this journey. Machine Learning can feel like magic, but I promise you it's just math, logic, and a lot of careful experimentation.

Since we are going from **absolute beginner to advanced**, I have designed a structured "Masterclass" syllabus for us.

Here is our roadmap:

- **Phase 1: The Intuition (Beginner)** - What is it, and how do we frame problems?
- **Phase 2: The Engine (Intermediate)** - How do models actually *learn*? (Loss & Optimization)
- **Phase 3: The Pitfalls (Intermediate)** - Overfitting, Evaluation, and the Bias-Variance tradeoff.
- **Phase 4: The Heavy Hitters (Advanced)** - Ensembles, SVMs, and Regularization.
- **Phase 5: The Real World (Expert)** - Messy data, ture engineering, and deployment.

Ask Qwen

Auto